

Structured Correlation Matrices in Stan

StanCon 2024

Sean Pinkney

sean.pinkney@gmail.com

Managing Director at Omnicom Media Group

Stan Developer



What is a correlation matrix?

- Symmetric
- 1s along the diagonal
- Off diagonal elements between $[-1, 1]$
- Positive semi-definite


That last requirement is the tough one.

An $n \times n$ matrix Σ is p.s.d iff

$$\mathbf{x}^\top \Sigma \mathbf{x} \geq 0 \quad \forall \mathbf{x} \in \mathbb{R}^n$$

Why is positive semi-definite tough?

Things get really, really constrained as the dimension of n increases



What is a structured correlation matrix?

Anything with structure.

“So any help you could give us would be most...helpful.” - Monthy Python

Types of structure

- Known values
- Block structure
- Bounds on correlation values

Cholesky factors

A positive definite matrix, Σ , can be factored into a lower triangular matrix, \mathbf{L} such that

$$\Sigma = \mathbf{L}\mathbf{L}^\top$$

Bonus:

A positive definite matrix, Σ , can also be factored into a strictly lower triangular matrix, \mathbf{L}_* and diagonal matrix, \mathbf{D} such that

$$\Sigma = \mathbf{L}_*\mathbf{D}\mathbf{L}_*^\top$$

Stan's current Cholesky factor of correlation matrix transform

```
1 data {
2   int<lower=0> K;
3 }
4 transformed data {
5   int K_choose_2 = choose(K, 2);
6 }
7 parameters {
8   // y is a vector K-choose-2 unconstrained parameters
9   vector[K_choose_2] y;
10 }
11 transformed parameters {
12   // L is a Cholesky factor of a K x K correlation matrix
13   cholesky_factor_corr[K] L = identity_matrix(K);
14   real log_det_jacobian = 0;
15   {
16     int counter = 1;
17     real sum_sqs;
18     vector[K_choose_2] z = tanh(y);
19     log_det_jacobian += sum(log1m(square(z)));
20
21     for (i in 2 : K) {
22       L[i, 1] = z[counter];
23       counter += 1;
24       sum_sqs = square(L[i, 1]);
25       for (j in 2 : (i - 1)) {
26         log_det_jacobian += 0.5 * log1m(sum_sqs);
27         L[i, j] = z[counter] * sqrt(1 - sum_sqs);
```

Motivating Example

Constrain the y “raw” parameters to be positive. This should make the correlation matrix all positive.

```
...  
    vector[K_choose_2] z = inv_logit(y); // z is between 0 and 1  
    log_det_jacobian += sum(log_inv_logit(y) + log1m_inv_logit(y));  
...
```


Let's first see how rejection sampling handles this

Sample a bunch and if **all** are positive then keep

```
using ArviZ,  
    Distributions,  
    PrettyTables  
  
xs = Matrix[];  
  
while length(xs) < 10000  
    X = rand(LKJ(5, 4))  
    if all(≥(0), X)  
        push!(xs, X)  
    end  
end
```

Mean should be about 0.28 and std 0.18

parameter	mean	std	hdi_2.5%	hdi_97.5%
String	Float64	Float64	Float64	Float64
x[1,1]	1.0	0.0	1.0	1.0
x[2,1]	0.282	0.181	0.0	0.611
x[3,1]	0.2806	0.1783	0.0001	0.6035
x[4,1]	0.2809	0.1815	0.0	0.6098
x[5,1]	0.2801	0.1808	0.0001	0.6115
x[1,2]	0.282	0.181	0.0	0.611
x[2,2]	1.0	0.0	1.0	1.0
x[3,2]	0.2819	0.1816	0.0	0.6119
x[4,2]	0.2787	0.1792	0.0	0.6091
x[5,2]	0.2827	0.1808	0.0	0.6109
x[1,3]	0.2806	0.1783	0.0001	0.6035
x[2,3]	0.2819	0.1816	0.0	0.6119
x[3,3]	1.0	0.0	1.0	1.0
x[4,3]	0.277	0.1799	0.0001	0.6068
x[5,3]	0.2792	0.1801	0.0004	0.6065
x[1,4]	0.2809	0.1815	0.0	0.6098
x[2,4]	0.2787	0.1792	0.0	0.6091
x[3,4]	0.277	0.1799	0.0001	0.6068
x[4,4]	1.0	0.0	1.0	1.0
x[5,4]	0.2778	0.1791	0.0003	0.6003
x[1,5]	0.2801	0.1808	0.0001	0.6115
x[2,5]	0.2827	0.1808	0.0	0.6109
x[3,5]	0.2792	0.1801	0.0004	0.6065

Output in Stan

```

1 # A tibble: 25 × 10
2   variable      mean median    sd   mad    q5   q95  rhat ess_bulk ess_tail
3   <chr>      <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>
4 1 Omega[1,1] 1      1      0      0      1      1    NA      NA      NA
5 2 Omega[2,1] 0.235  0.206  0.167  0.175  0.0195  0.549  1.00  44411.  20215.
6 3 Omega[3,1] 0.235  0.205  0.168  0.177  0.0195  0.555  1.00  51831.  22326.
7 4 Omega[4,1] 0.235  0.206  0.167  0.177  0.0200  0.555  1.00  46224.  21451.
8 5 Omega[5,1] 0.235  0.206  0.166  0.175  0.0198  0.551  1.00  50078.  21105.
9 6 Omega[1,2] 0.235  0.206  0.167  0.175  0.0195  0.549  1.00  44411.  20215.
10 7 Omega[2,2] 1      1      0      0      1      1    NA      NA      NA
11 8 Omega[3,2] 0.280  0.257  0.165  0.176  0.0510  0.587  1.00  53901.  27702.
12 9 Omega[4,2] 0.280  0.258  0.167  0.178  0.0497  0.588  1.00  45097.  25476.
13 10 Omega[5,2] 0.280  0.258  0.166  0.175  0.0498  0.588  1.00  47192.  24449.
14 11 Omega[1,3] 0.235  0.205  0.168  0.177  0.0195  0.555  1.00  51831.  22326.
15 12 Omega[2,3] 0.280  0.257  0.165  0.176  0.0510  0.587  1.00  53901.  27702.
16 13 Omega[3,3] 1      1      0      0      1      1    NA      NA      NA
17 14 Omega[4,3] 0.325  0.307  0.165  0.176  0.0870  0.624  1.00  51438.  29513.
18 15 Omega[5,3] 0.324  0.307  0.164  0.175  0.0870  0.621  1.00  48371.  28163.
19 16 Omega[1,4] 0.235  0.206  0.167  0.177  0.0200  0.555  1.00  46224.  21451.
20 17 Omega[2,4] 0.280  0.258  0.167  0.178  0.0497  0.588  1.00  45097.  25476.
21 18 Omega[3,4] 0.325  0.307  0.165  0.176  0.0870  0.624  1.00  51438.  29513.
22 19 Omega[4,4] 1      1      0      0      1      1    NA      NA      NA
23 20 Omega[5,4] 0.368  0.355  0.163  0.175  0.122   0.657  1.00  50304.  30427.
24 21 Omega[1,5] 0.235  0.206  0.166  0.175  0.0198  0.551  1.00  50078.  21105.
25 22 Omega[2,5] 0.280  0.258  0.166  0.175  0.0498  0.588  1.00  47192.  24449.
26 23 Omega[3,5] 0.324  0.307  0.164  0.175  0.0870  0.621  1.00  48371.  28163.
27 24 Omega[4,5] 0.368  0.355  0.163  0.175  0.122   0.657  1.00  50304.  30427.

```

Wtf?

???

A foray into Cholesky factors



Andre Cholesky

Cholesky factors

The value $C_{i,j} \in \mathbf{C}$ is a correlation value that is between $-1, 1$

$$\mathbf{L} = \begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ C_{2,1} & \sqrt{1 - L_{2,1}^2} & 0 & \cdots & 0 \\ C_{3,1} & (C_{3,2} - L_{3,1}L_{2,1}) / L_{2,2} & \sqrt{1 - L_{3,1}^2 - L_{3,2}^2} & \cdots & 0 \\ \vdots & \vdots & \vdots & \cdots & \vdots \\ C_{n,1} & (C_{n,2} - L_{n,1}L_{2,1}) / L_{2,2} & (C_{n,3} - L_{n,1}L_{2,1} - L_{n,2}L_{3,2}) / L_{3,3} & \cdots & \sqrt{1 - \sum_{k=1}^{n-1} L_{n,k}^2} \end{pmatrix}$$

Let's try something

Let's condense it into an equation

$$L_{j,j} = \sqrt{1 - \sum_{k=1}^{j-1} L_{j,k}^2}$$
$$L_{i,j} = \frac{1}{L_{j,j}} \left(C_{i,j} - \sum_{k=1}^{j-1} L_{i,k} L_{j,k} \right) \quad \text{for } i > j$$

Re-arrange and add in user specified **lower** and **upper** bounds

$$-L_{j,j} + \sum_{k=1}^{j-1} L_{i,k} L_{j,k} \leq \mathbf{a_{i,j}} < C_{i,j} < \mathbf{b_{i,j}} \leq L_{j,j} + \sum_{k=1}^{j-1} L_{i,k} L_{j,k}$$

Continuing

$$\begin{aligned}
 & -\sqrt{1 - \sum_{k=1}^{j-1} L_{i,k}^2} \leq \frac{\mathbf{a}_{i,j} - \sum_{k=1}^{j-1} L_{i,k} L_{j,k}}{L_{j,j}} < L_{i,j} < \frac{\mathbf{b}_{i,j} - \sum_{k=1}^{j-1} L_{i,k} L_{j,k}}{L_{j,j}} \leq \sqrt{1 - \sum_{k=1}^{j-1} L_{i,k}^2} \\
 & \max \left\{ -\sqrt{1 - \sum_{k=1}^{j-1} L_{j,k}^2}, \frac{\mathbf{a}_{i,j} - \sum_{k=1}^{j-1} L_{i,k} L_{j,k}}{L_{j,j}} \right\} < L_{i,j} < \min \left\{ \sqrt{1 - \sum_{k=1}^{j-1} L_{j,k}^2}, \frac{\mathbf{b}_{i,j} - \sum_{k=1}^{j-1} L_{i,k} L_{j,k}}{L_{j,j}} \right\}
 \end{aligned}$$

This gives us bounds but we need to be really careful

The Jacobian is “easy”

Extra Care

Assume we have 3×3 matrix that we wish to constrain all correlation values to be negative.

Let us choose $C_{2,1} = C_{3,1} = \frac{-1}{\sqrt{2}}$ and solve for the bounds of $C_{3,2}$:

$$\begin{aligned}\max \left\{ -1, -\sqrt{1 - C_{3,1}^2} \right\} &< \frac{B - C_{2,1}C_{3,1}}{L_{2,2}} < \min \left\{ 0, \sqrt{1 - C_{3,1}^2} \right\} \\ -\sqrt{0.5} &< \frac{B - 0.5}{\sqrt{0.5}} < 0 \\ \implies 0 &< B < 0.5\end{aligned}$$

Note

This is all in Pinkney (2024)

The Stan code

```
1 matrix cholesky_corr_constrain_lp (vector col_one_raw, vector off_raw,
2                                   real lb, real ub) {
3   int K = num_elements(col_one_raw) + 1;
4   matrix[K, K] L = rep_matrix(0, K, K);
5   L[1, 1] = 1;
6   L[2, 1] = lb_ub_lp(col_one_raw[1], lb, ub);
7   L[2, 2] = sqrt(1 - L[2, 1]^2);
8
9   int cnt = 1;
10
11   for (i in 3:K) {
12     L[i, 1] = lb_ub_lp(col_one_raw[2:K - 1], lb, ub);
13     real l_ij_old = loglm(L[i, 1]^2);
14     for (j in 2:i - 1) {
15       real b1 = dot_product(L[j, 1:(j - 1)], L[i, 1:(j - 1)]);
16       real stick_length = exp(0.5 * l_ij_old);
17       real low = max({-stick_length, (lb - b1) / L[j, j]});
18       real up = min({stick_length, (ub - b1) / L[j, j]});
19       L[i, j] = lb_ub_lp(off_raw[cnt], low, up);
20       l_ij_old = log_diff_exp(l_ij_old, 2 * log(abs(L[i, j])));
21       cnt += 1;
22     }
23     L[i, i] = exp(0.5 * l_ij_old);
24   }
25   return L;
26 }
```

This is a bit more stable than what's in the paper but also more opaque.

The `exp()`'s here are just calculating stuff on the log-scale.

Line 13 I take the log and try to stick with this as much as possible.

Back to the motivating example

The new method that gives the right answer!

```
1 # A tibble: 25 × 10
2   variable    mean median    sd   mad    q5   q95  rhat ess_bulk ess_tail
3   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>
4 1 Omega[1,1] 1      1      0      0      1      1    NA        NA        NA
5 2 Omega[2,1] 0.278  0.256 0.179 0.198 0.0281 0.607 1.00    29806.    17895.
6 3 Omega[3,1] 0.279  0.258 0.180 0.200 0.0273 0.609 1.00    34574.    18758.
7 4 Omega[4,1] 0.279  0.258 0.180 0.199 0.0276 0.607 1.00    29344.    16380.
8 5 Omega[5,1] 0.281  0.259 0.181 0.201 0.0291 0.610 1.00    36064.    18496.
9 6 Omega[1,2] 0.278  0.256 0.179 0.198 0.0281 0.607 1.00    29806.    17895.
10 7 Omega[2,2] 1      1      0      0      1      1    NA        NA        NA
11 8 Omega[3,2] 0.279  0.258 0.179 0.200 0.0283 0.605 1.00    37118.    21150.
12 ...
```

What else can this do?

- Known values
- Specific bounds on certain values
- Block structure

Let's do a really motivating example

An even more motivating example

We want a 6 x 6 correlation matrix Σ where

$$\sigma_{2,1} = 0, \sigma_{4,3} = 0$$

$$\sigma_{4,1} \in (-0.8, 0), \sigma_{5,3} \in (0.3, 0.7)$$

```
1  --
2  0.  --
3  NA  NA  --
4  (-0.8, 0) NA  0  --
5  NA  NA  (0.3, 0.7) NA  --
6  NA  NA  NA  NA  NA  --
```

Output

```

1 variable      mean median      sd      mad      q5      q95      rhat ess_bulk ess_tail
2 <chr>          <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>    <dbl>    <dbl>
3 Omega[1,1]    1.0000   1.0000  0.0000  0.0000  1.0000  1.0000  NA        NA        NA
4 Omega[2,1]    0.0000   0.0000  0.0000  0.0000  0.0000  0.0000  NA        NA        NA
5 Omega[3,1]   -0.0056  -0.0101  0.2484  0.2559 -0.4211  0.4068  0.9999  4410.0862 3220.8831
6 Omega[4,1]   -0.2124  -0.1881  0.1512  0.1579 -0.5044 -0.0173  1.0003  4729.1236 2149.4316
7 Omega[5,1]   -0.0057  -0.0065  0.2621  0.2772 -0.4435  0.4217  1.0000  4657.2520 3017.0624
8 Omega[6,1]   -0.0018  -0.0002  0.2553  0.2599 -0.4284  0.4143  1.0013  4674.8017 2903.7293
9
10 Omega[2,2]   1.0000   1.0000  0.0000  0.0000  1.0000  1.0000  NA        NA        NA
11 Omega[3,2]  -0.0043  -0.0078  0.2499  0.2636 -0.4178  0.4093  1.0001  4023.8517 2705.3045
12 Omega[4,2]   0.0043   0.0027  0.2594  0.2767 -0.4159  0.4328  1.0007  5443.1496 3098.8828
13 Omega[5,2]   0.0014   0.0093  0.2575  0.2682 -0.4298  0.4201  1.0017  4128.4046 2987.4422
14 Omega[6,2]   0.0058   0.0064  0.2538  0.2693 -0.4197  0.4224  1.0006  4736.2346 2668.6795
15
16 Omega[3,3]   1.0000   1.0000  0.0000  0.0000  1.0000  1.0000  NA        NA        NA
17 Omega[4,3]  -0.0000   0.0000  0.0000  0.0000  0.0000  0.0000  1.0014  4280.0269  NA
18 Omega[5,3]   0.4128   0.3934  0.0882  0.0884  0.3072  0.5888  1.0012  4809.9923 2377.6445
19 Omega[6,3]  -0.0032  -0.0064  0.2624  0.2747 -0.4394  0.4265  1.0022  4673.5773 3037.9778
20
21 Omega[4,4]   1.0000   1.0000  0.0000  0.0000  1.0000  1.0000  NA        NA        NA
22 Omega[5,4]   0.0024   0.0020  0.2371  0.2475 -0.3873  0.3970  1.0002  3895.9938 2933.7159
23 Omega[6,4]  -0.0004  -0.0033  0.2576  0.2659 -0.4183  0.4327  1.0019  4303.5343 2949.4916
24
25 Omega[5,5]   1.0000   1.0000  0.0000  0.0000  1.0000  1.0000  NA        NA        NA
26 Omega[6,5]   0.0009  -0.0018  0.2533  0.2675 -0.4113  0.4213  1.0002  4139.6808 2995.0615

```

Do you want to see the Stan code for this?

```
1  matrix cholesky_corr_constrain_lp(int K, vector raw, int N_blocks,
2                                  array[,] int res_index,
3                                  array[] int res_id,
4                                  array[,] int known_index,
5                                  vector known_vals,
6                                  vector lb, vector ub) {
7    matrix[K, K] L = rep_matrix(0, K, K);
8    int cnt = 1;
9    int N_res = num_elements(res_id);
10   int N_known = size(known_vals);
11   vector[N_blocks] x_cache;
12   array[N_blocks] int res_id_cnt = ones_int_array(N_blocks);
13   int res_row = 1;
14   int known_row = 1;
15
16   L[1, 1] = 1;
17   if (known_index[known_row, 1] == 2) {
18     L[2, 1] = known_vals[1];
19     known_row += 1;
20   } else {
21     L[2, 1] = lb_ub_lp(raw[cnt], lb[cnt], ub[cnt]);
22     if (res_index[res_row, 1] == 2) {
23       x_cache[res_id[res_row]] = L[2, 1];
24       res_id_cnt[res_id[res_row]] += 1;
25       res_row += 1;
26     }
27     cnt += 1;
```

The Future

LDL parameterization is done. It's even more stable.

Look-ahead given bound constraints and known values to ensure the impossible bounds don't hit

Explore new priors for correlation matrices that exploit structure

References:

Pinkney, Sean. 2024. "A Short Note on a Flexible Cholesky Parameterization of Correlation Matrices."
<https://arxiv.org/abs/2405.07286>.